

Altamash Azam

AI Engineer (Trainee) – Agentic Systems & LLM Applications
Kolkata, WB, India · altamash76azam@outlook.com · +91 7643946866
linkedin.com/in/altamash-azam-757120422 · github.com/Altamashazam7

SUMMARY

AI Engineer trainee focused on agentic LLM systems and retrieval-augmented generation. Builds autonomous agent workflows with LangChain/LangGraph, RAG pipelines over vector and relational data, and evaluates them against explicit baselines rather than anecdotal testing. Comfortable owning a project from prompt design through to a working, demoable prototype.

TECHNICAL SKILLS

Core (hands-on, can go deep): Python, LangChain, LangGraph, RAG pipeline design, Prompt Engineering, Vector Search (FAISS, ChromaDB)

Also used: CrewAI, AutoGen, Pinecone, PostgreSQL, MongoDB, PyTorch, scikit-learn

Models / APIs: GPT-4, Claude 3.5, Gemini 1.5 (API integration, prompt routing, multimodal input handling)

Tools & Infra: Git, Docker, AWS (EC2, S3), GCP, CI/CD basics, VS Code

EXPERIENCE & INTERNSHIP

Data Analytics Intern

Jan 2025 – Feb 2025

DivSriVin Technologies (Ref ID: DIVSR1004)

Chennai, TN, India (Remote)

- Completed an industry-mentored internship in Data Analytics and Data Science under CEO/Mentor Mr. Vishnuu A.
- Built a K-Nearest Neighbors (KNN) classification model in `scikit-learn` to predict subscriber churn as the capstone project (*Dish TV Churn Prediction*).
- Engineered features from ~45,000 subscriber records based on account tenure, email verification status, and network activity flags; capstone was reviewed, verified, and officially approved upon completion.

KEY PROJECTS

Document Q&A Engine with RAG Pipeline *Python | FastAPI | LangChain | FAISS | OpenAI / Hugging Face*

- Engineered an in-memory Retrieval-Augmented Generation (RAG) pipeline to perform vector similarity search over uploaded unstructured PDF documents.
- Implemented sliding-window text chunking strategies (500 tokens with overlapping windows) and FAISS vector indexing to optimize context retrieval accuracy and reduce prompt token bloat by 40%.
- Built a FastAPI backend streaming responses with grounded source citations to eliminate LLM hallucination risks.

Distributed API Rate Limiter & Redis Caching Layer

Node.js | Express | Redis | Docker

- Implemented sliding-window rate limiting middleware using Redis sorted sets (ZSET) to protect public endpoints from API abuse and DoS threats.
- Built an in-memory response caching layer for high-throughput GET routes, dropping redundant database query loads by 60% and cutting latency during peak traffic.
- Containerized the application service alongside Redis using Docker Compose for simple local orchestration and consistent deployments.

EDUCATION

B.Tech in Computer Science & Engineering

2023 – 2027 (Expected)

Budge Budge Institute of Technology

Relevant coursework: Artificial Intelligence, Machine Learning, Data Structures, Database Systems.

CERTIFICATIONS

- Data Analytics & Data Science Internship Certificate** – DivSriVin Technologies (Ref ID: DIVSR1004)
- Machine Learning Specialization** – Andrew Ng, DeepLearning.AI / Stanford (Coursera)
- AWS Cloud Practitioner** – Amazon Web Services
- Deep Learning with PyTorch** – Udemy
- Python Programming** – Udemy